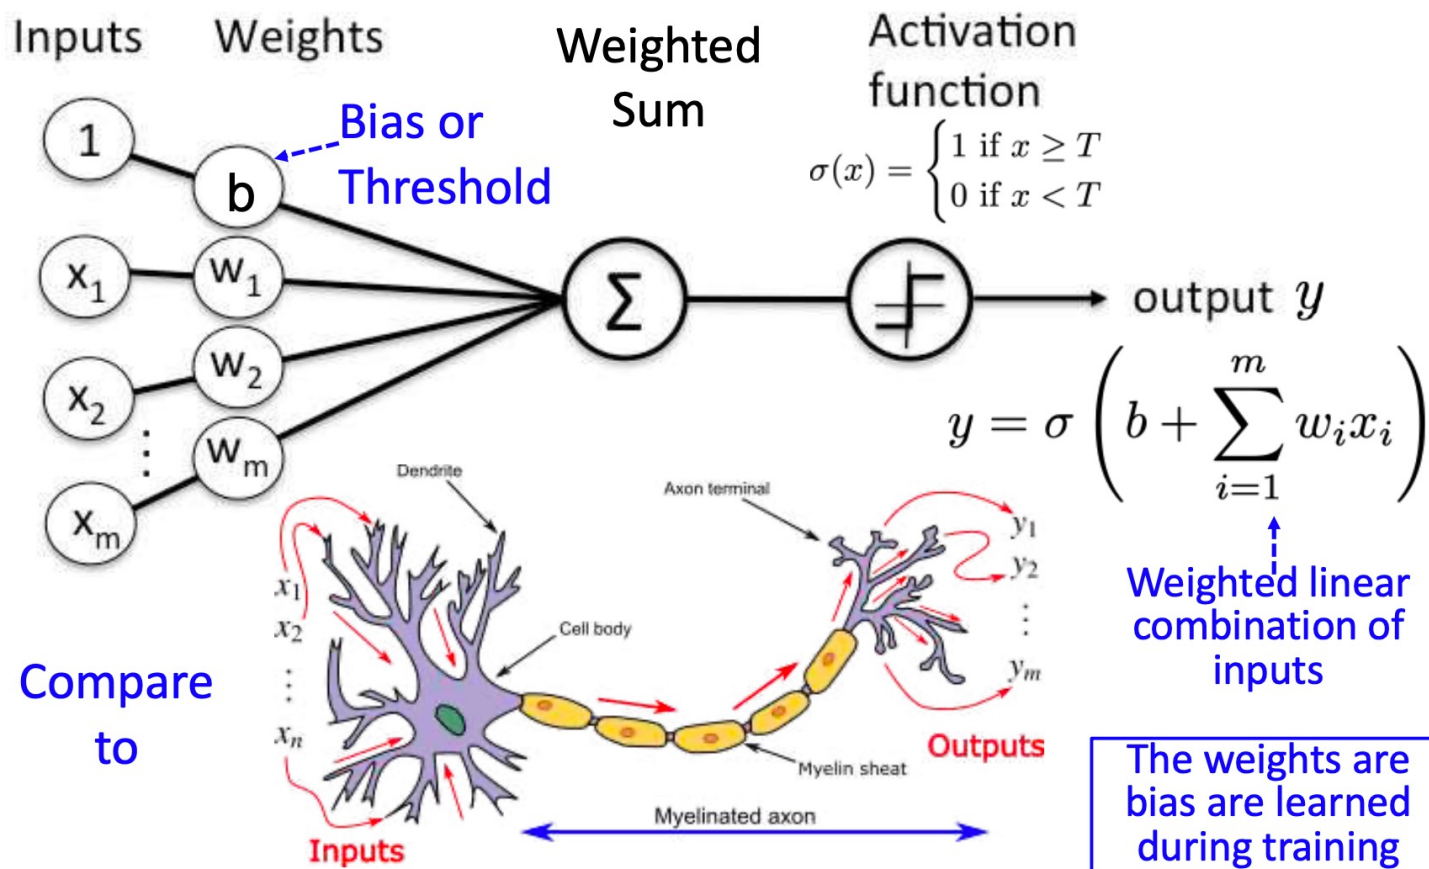


Deep Differentiable Logic Gate Networks

Felix Petersen, Christian Borgelt, Hilde Kuehne, Oliver Deussen

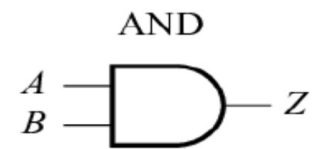
Interpreted by:
John Tan Chong Min

Perceptron



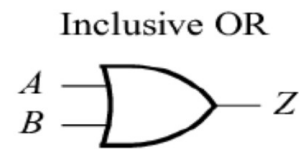
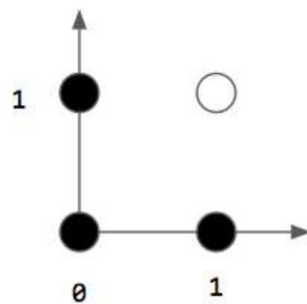
Source: Mehul Motani's Neural Network Notes

Logic Gates



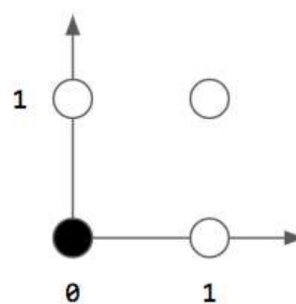
Inputs		Output
A	B	Z
0	0	0
0	1	0
1	0	0
1	1	1

AND



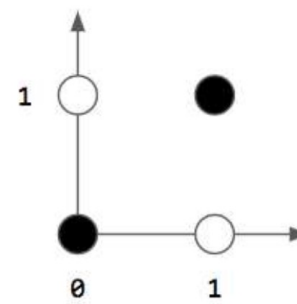
Inputs		Output
A	B	Z
0	0	0
0	1	1
1	0	1
1	1	1

OR

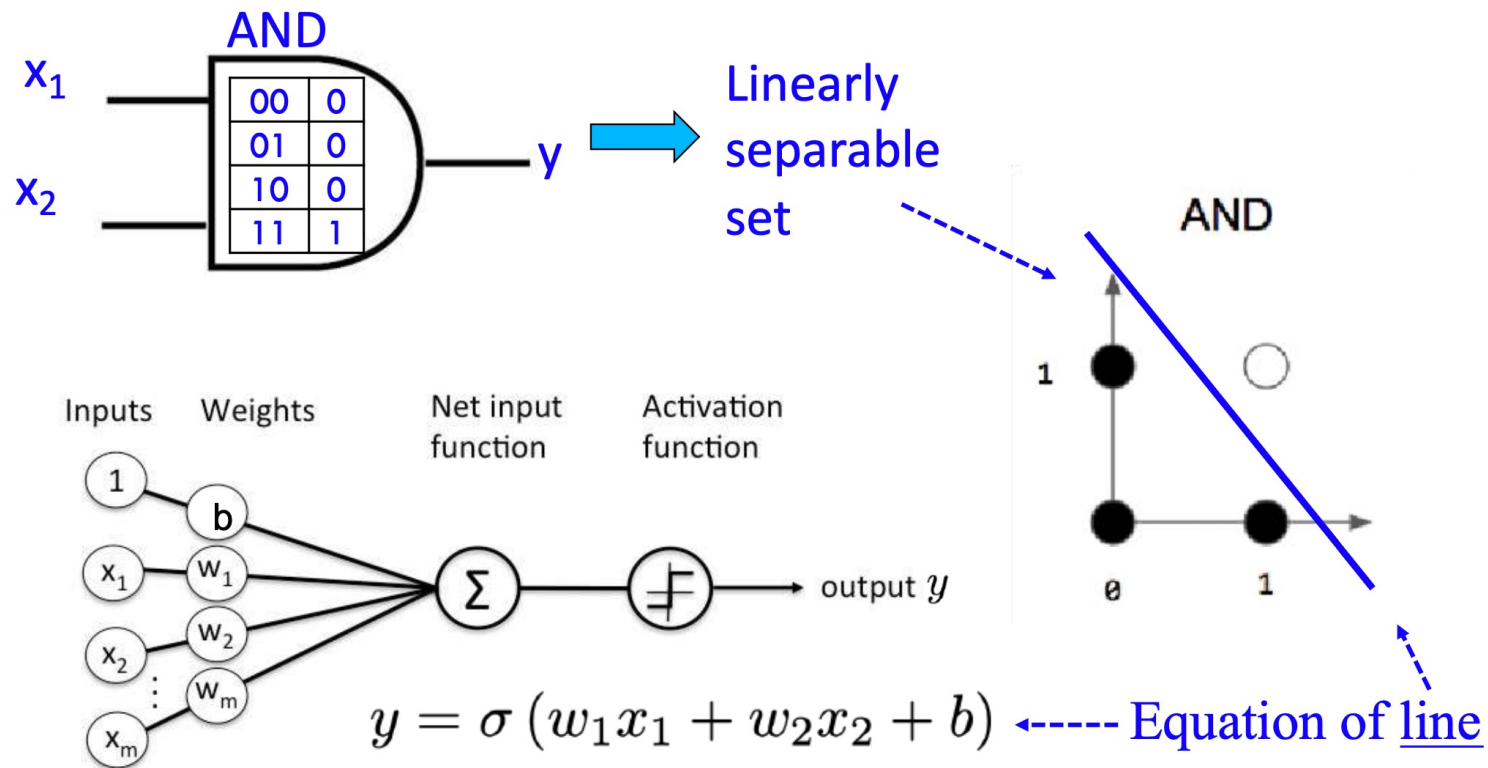


Inputs		Output
A	B	Z
0	0	0
0	1	1
1	0	1
1	1	0

XOR

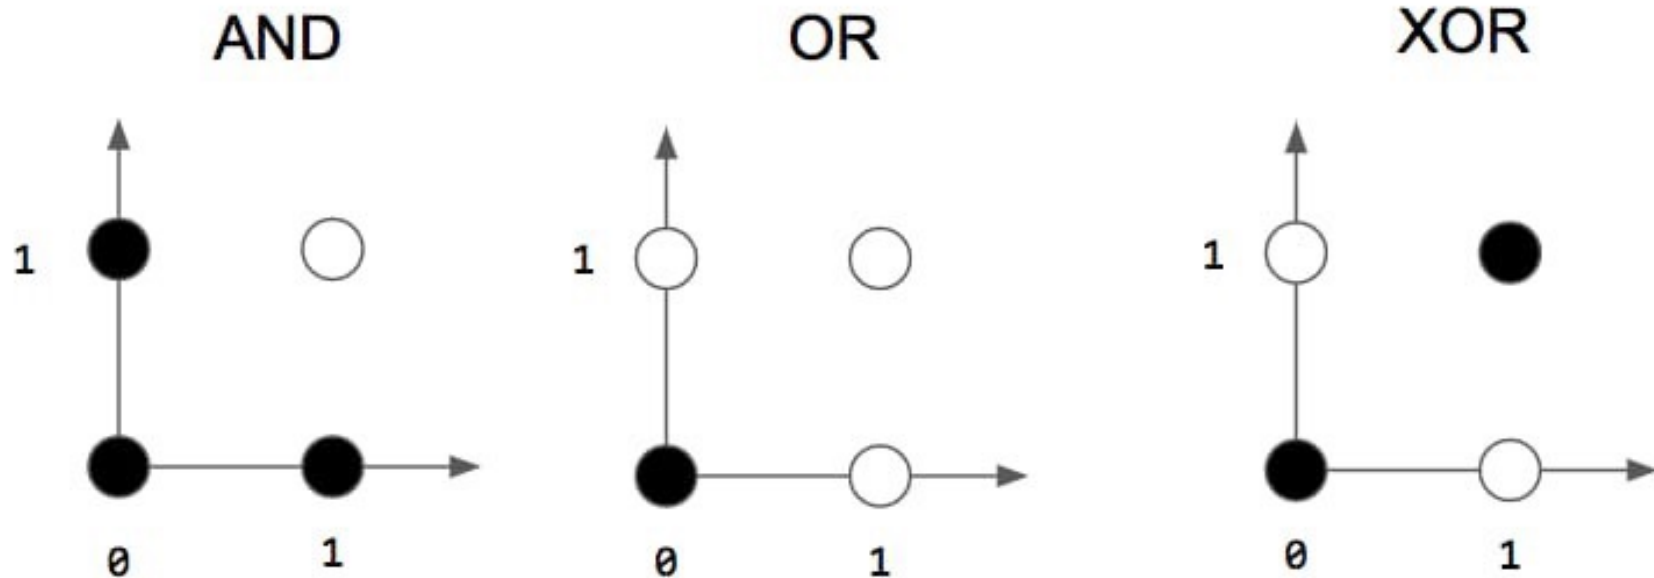


Perceptron is a linear separator

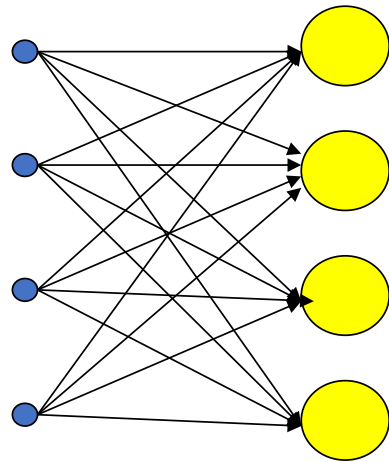


The XOR problem

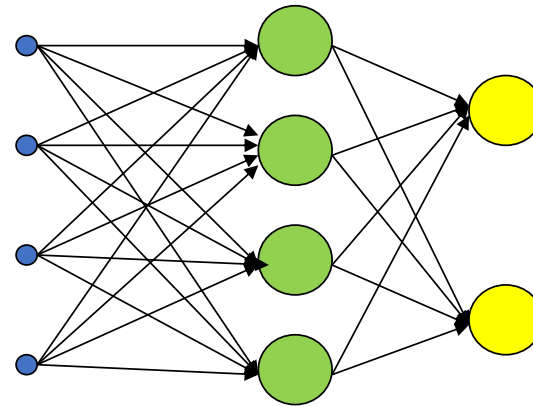
- Try to draw a line to separate the black and white circles



Multi-Layer Perceptron

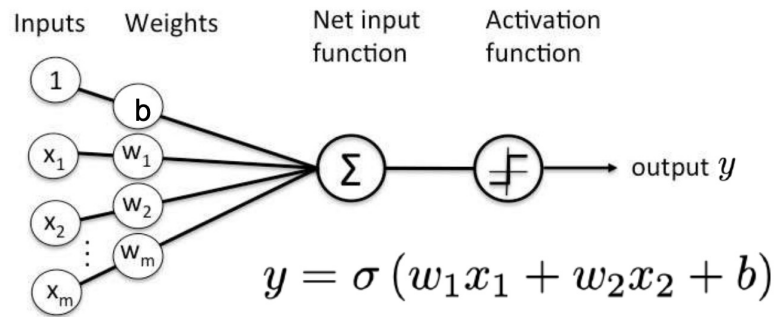


Single Layer
(Perceptron)



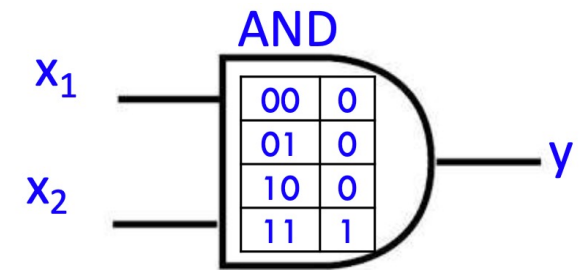
Multiple Layers
(Multi-Layer Perceptron)

Logic Gates vs Neurons

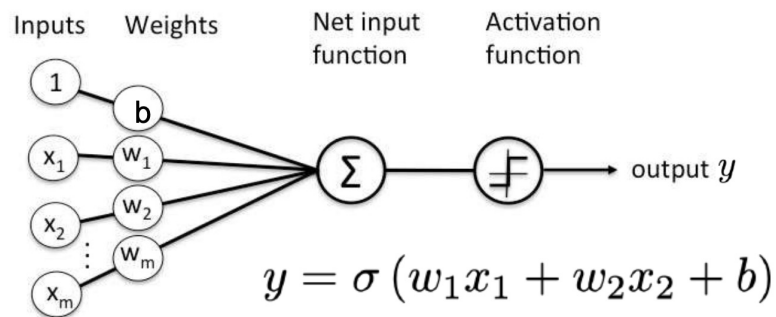


specialize →

← need more to emulate

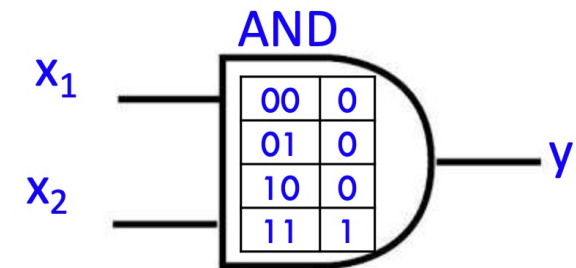


Logic Gates vs Neurons



specialize \rightarrow

\leftarrow need more to emulate



Only models linear decision boundaries!

Neural-like: Flexible learning of function

Slow learning and inference

Can have non-linear decision boundaries!
E.g. XOR

Algorithm-like: Fixed function

Fast learning and inference

What Logic Gates to model?

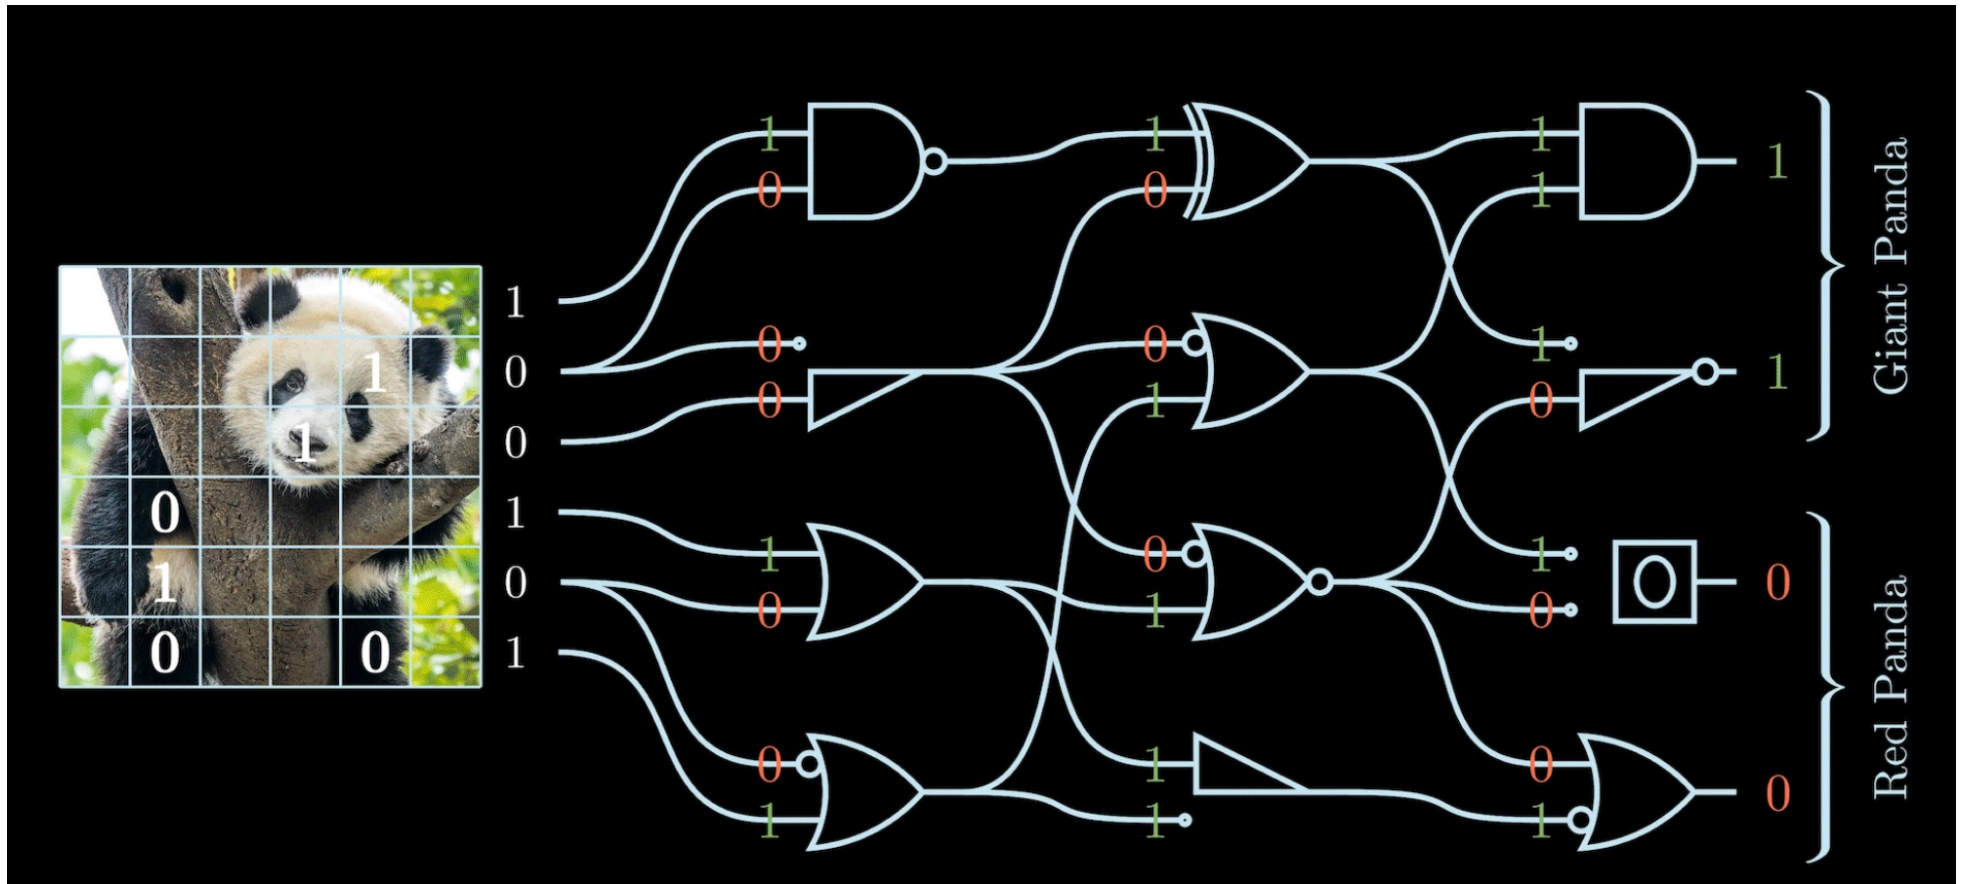
- Model every possible 2-input gate



Table 1: List of all real-valued binary logic ops.

ID	Operator	real-valued	00	01	10	11
0	False	0	0	0	0	0
1	$A \wedge B$	$A \cdot B$	0	0	0	1
2	$\neg(A \Rightarrow B)$	$A - AB$	0	0	1	0
3	A	A	0	0	1	1
4	$\neg(A \Leftarrow B)$	$B - AB$	0	1	0	0
5	B	B	0	1	0	1
6	$A \oplus B$	$A + B - 2AB$	0	1	1	0
7	$A \vee B$	$A + B - AB$	0	1	1	1
8	$\neg(A \vee B)$	$1 - (A + B - AB)$	1	0	0	0
9	$\neg(A \oplus B)$	$1 - (A + B - 2AB)$	1	0	0	1
10	$\neg B$	$1 - B$	1	0	1	0
11	$A \Leftarrow B$	$1 - B + AB$	1	0	1	1
12	$\neg A$	$1 - A$	1	1	0	0
13	$A \Rightarrow B$	$1 - A + AB$	1	1	0	1
14	$\neg(A \wedge B)$	$1 - AB$	1	1	1	0
15	True	1	1	1	1	1

Classification

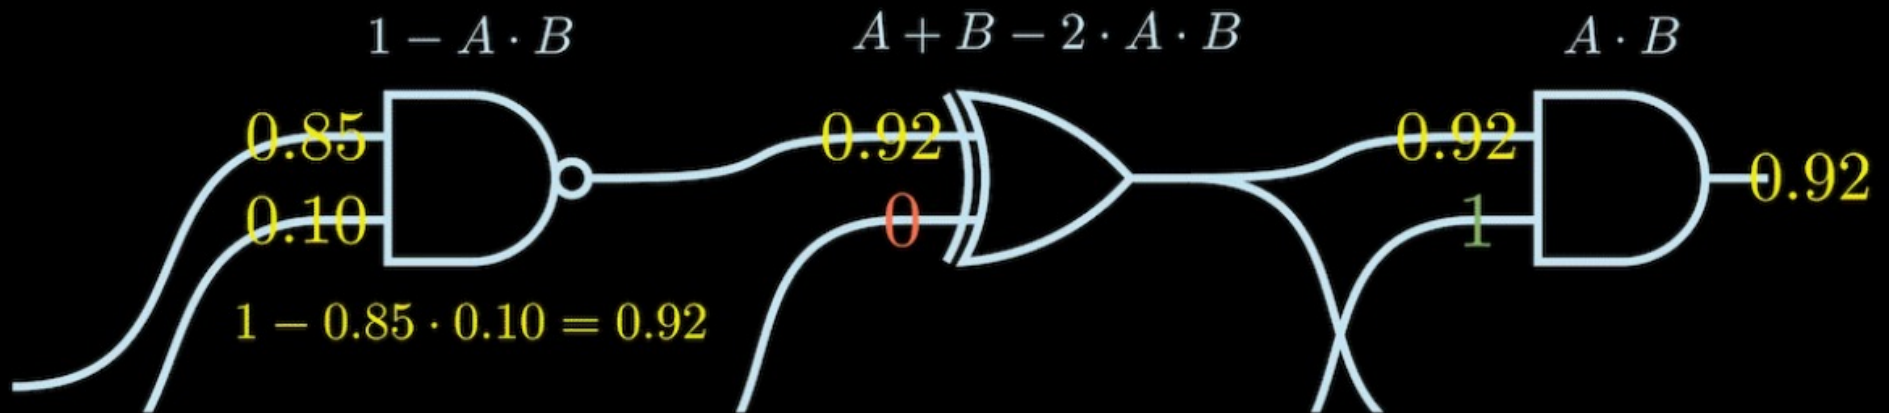


Difficulty of training logic gates

- Logic gates perform one fixed function, inherently not differentiable
- Could use evolutionary means to learn the logic gates -> not scalable
- Make network differentiable:
 - Link inputs to outputs using a differentiable function
 - Use some probabilistic way to select logic gates

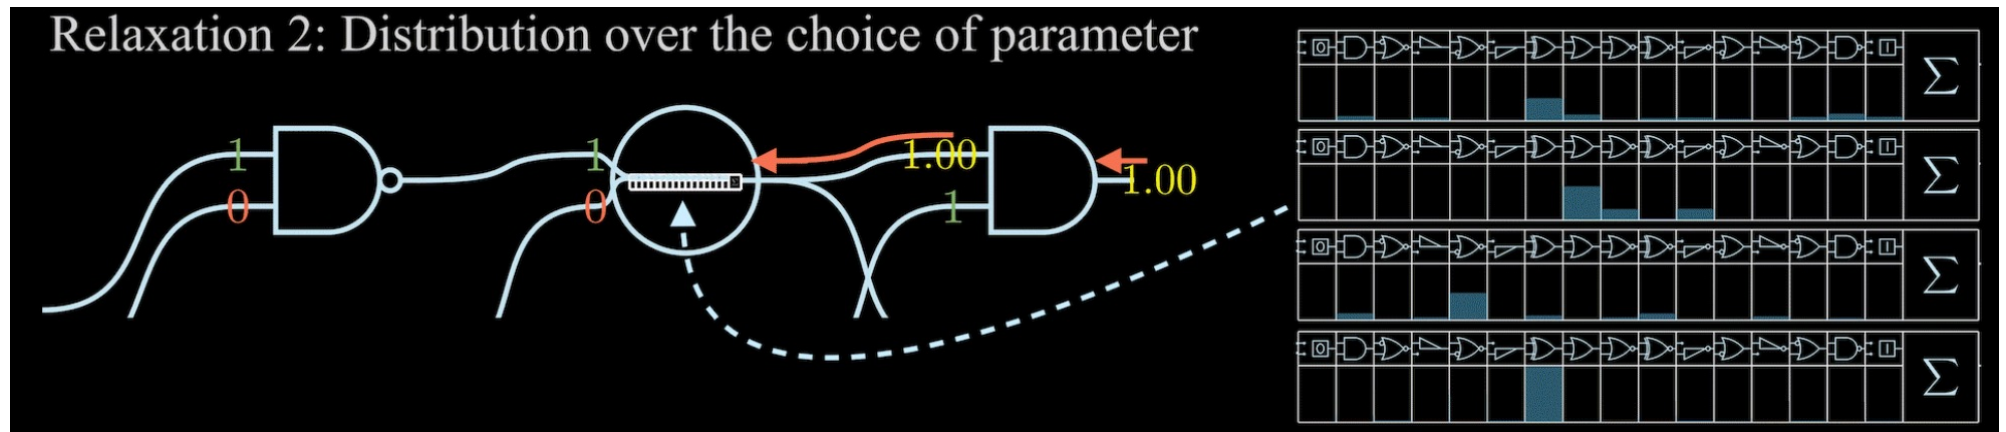
Make output a function of the inputs

Relaxation 1: Real-valued Logics



$$a' = f_i(a_1, a_2)$$

Make the gate selectable by backprop



$$a' = \sum_{i=0}^{15} \boxed{p_i} \cdot f_i(a_1, a_2) = \sum_{i=0}^{15} \boxed{\frac{e^{w_i}}{\sum_j e^{w_j}}} \cdot f_i(a_1, a_2)$$

Categorical Distribution via Softmax

Training Setup

- **Initialization:** Randomly initialize connections and parameterization of each neuron (normal distribution), Adam optimizer, 0.01 learning rate
- **Logic Gate selection:** Each neuron will be converted to the logic gate with the highest probability based on the categorical probability distribution

Configuring to Output (Classification)

- k classes, n output neurons. Group neuron outputs in groups of size n/k , then aggregate the classification scores. Train by passing through all outputs with a softmax, and class with largest sum of post-softmax values is winner

Aggregation of Output Neurons Now, we may have n output neurons $a_1, a_2, \dots, a_n \in [0, 1]$, but we may want the logic gate network to only predict $k < n$ values of a larger range than $[0, 1]$. Further, we may want to be able to produce graded outputs. Thus, we can aggregate the outputs as

$$\hat{y}_i = \sum_{j=i \cdot n/k + 1}^{(i+1) \cdot n/k} a_j / \tau + \beta \quad (3)$$

where τ is a normalization temperature and β is an optional offset.

Results (Tabular - Classification)

Table 3: Results for the Adult and Breast Cancer data sets averaged over 10 runs.

Adult	Acc.	# Param.	Infer. Time	Space
Decision Tree Learner	79.5%	≈ 50	86ns	$\approx 130\text{B}$
Logistic Regression	84.8%	234	63ns	936B
Neural Network	84.9%	3810	635ns	15KB
Diff Logic Net (<i>ours</i>)	84.8%	1280	5.1ns	640B
Breast Cancer	Acc.	# Param.	Infer. Time	Space
Decision Tree Learner	71.9%	≈ 100	82ns	$\approx 230\text{B}$
Logistic Regression	72.9%	104	34ns	416B
Neural Network	75.3%	434	130ns	1.4KB
Diff Logic Net (<i>ours</i>)	76.1%	640	2.8ns	320B

Fast inference, good accuracy for classification tasks in tabular datasets

Results (Binarized MNIST)

Table 4: Results for MNIST, all of our results are averaged over 10 runs. Times (T.) are inference times per image, the GPU is an NVIDIA A6000, and the CPU is a single thread at 2.5 GHz. For our experiments, i.e., the top block, we use binarized MNIST.

MNIST	Acc.	# Param.	Space	T. [CPU]	T. [GPU]	OPs	FLOPs
Linear Regression	91.6%	4 010	16KB	$3\mu s$	2.4ns	(4M)	4K
Neural Network (<i>small</i>)	97.92%	118 282	462KB	$14\mu s$	12.4ns	(236M)	236K
Neural Network	98.40%	22 609 930	86MB	2.2ms	819ns	(45G)	45M
Diff Logic Net (<i>small</i>)	97.69%	48 000	23KB	625ns	6.3ns	48K	—
Diff Logic Net	98.47%	384 000	188KB	$7\mu s$	(50ns)	384K	—

Fast inference, good accuracy for larger image classification

Results (CIFAR)

Table 5: Results on CIFAR-10. Times (T.) are inference times per image, the GPU is an NVIDIA A6000, and the CPU is a single thread at 2.5 GHz. For our experiments, i.e., the top block, we use a color-channel resolution of 4 for the first 3 lines and a color-channel resolution of 32 for the *large* models. The other baselines were provided with the full resolution of 256 color-channel values. The numbers in parentheses are extrapolated / estimated.

CIFAR-10	Acc.	# Param	Space	T. [CPU]	T. [GPU]	OPs	FLOPs
Neural Network (color-ch. res. = 4)	50.79%	12.6M	48MB	1.2ms	370ns	(25G)	25M
Diff Logic Net (<i>small</i>)	51.27%	48K	24KB	1.3 μ s	19ns	48K	—
Diff Logic Net (<i>medium</i>)	57.39%	512K	250KB	7.3 μ s	29ns	512K	—
Diff Logic Net (<i>large</i>)	60.78%	1.28M	625KB	(18 μ s)	(73ns)	1.28M	—
Diff Logic Net (<i>large</i> × 2)	61.41%	2.56M	1.22MB	(37 μ s)	(145ns)	2.56M	—
Diff Logic Net (<i>large</i> × 4)	62.14%	5.12M	2.44MB	(73 μ s)	(290ns)	5.12M	—
<i>Best Fully-Connected Baselines (color-ch. res. = 256)</i>							
Regularized SReLU NN [28]	68.70%	20.3M	77MB	1.9ms	565ns	(40G)	40M
Student-Teacher NN [52]	65.8%	1M	4MB	112 μ s	243ns	(2G)	2M
Student-Teacher NN [52]	74.3%	31.6M	121MB	2.9ms	960ns	(63G)	63M

Fast inference, moderate accuracy for larger image classification

Exponential Growth of Gates!

Table 6: Logic gate network architectures.

Dataset	Model	Layers	Neurons / layer	Total num. of p.		τ
MONK-1	—	6	24	144		1
MONK-2	—	6	12	72		1
MONK-3	—	6	12	72		1
Adult	—	5	256	32	1 280	1/0.075
Breast Cancer	—	5	128	8	640	1/0.1
MNIST	small	6	8 000	128	48 000	1/0.1
	normal	6	64 000	2048	384 000	1/0.03
CIFAR-10	small	4	12 000	1024	48 000	1/0.03
	medium	4	128 000		512 000	1/0.01
	large	5	256 000		1 280 000	1/0.01
	large×2	5	512 000		2 560 000	1/0.01
	large×4	5	1 024 000		5 120 000	1/0.01

Exponential
Increase

Requires more
neurons
than MLP
(in orange)



Limitations

- Expensive Training – Higher training cost compared to conventional neural networks. Mainly due to the differentiable categorical distribution used to decide the logic gate
- Convolutions and Residual connections are not modelled
- Limited to small architectures due to exponential scaling of number of logic gates in subsequent layers

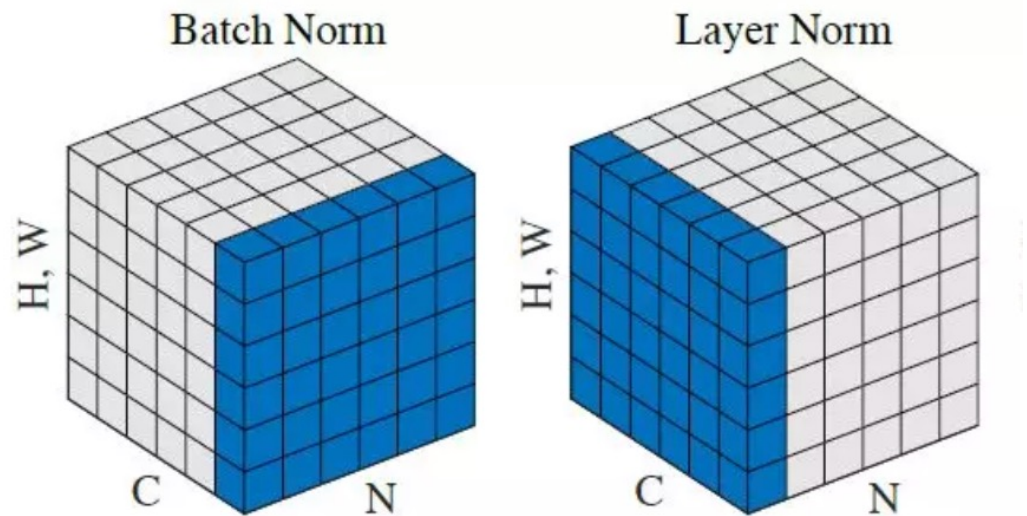
Some of my thoughts

Benefits and drawbacks (my own thoughts)

- Benefits:
 - No gradient vanishing / exploding
 - Mimics biological neurons better – output is always limited to 1
- Drawbacks:
 - Requires exponential amount of logic gates to model the input-output relation
 - Requires a lot of logic gate nodes as compared to the normal neurons
- Maybe take the paper's idea of constraining outputs, but use a more expressive architecture?

Artificial Tools used in Neural Networks to counter vanishing/exploding gradients

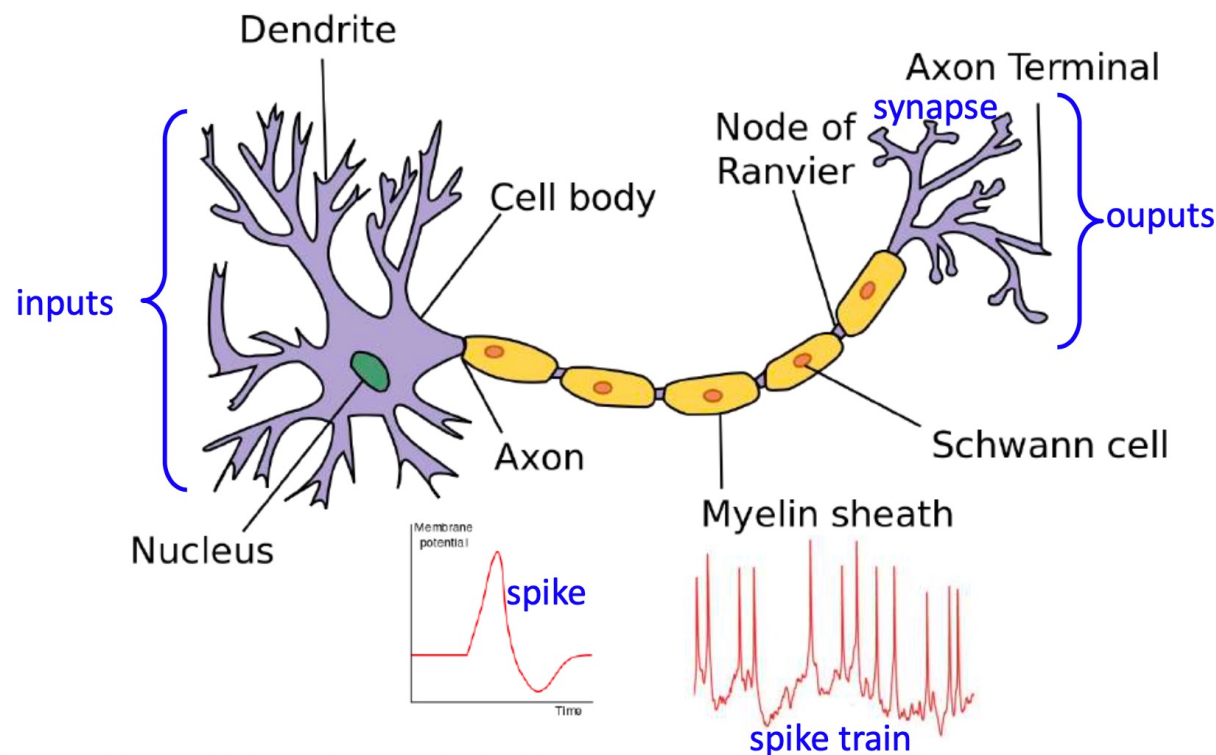
- Scaling of inputs to zero mean and unit variance (e.g. ImageNet)
- Normalization
 - Batch Norm
 - Layer Norm
- Gated RNNs
 - Forget Gate
 - Input Gate
 - Output Gate



<https://paperswithcode.com/method/layer-normalization>

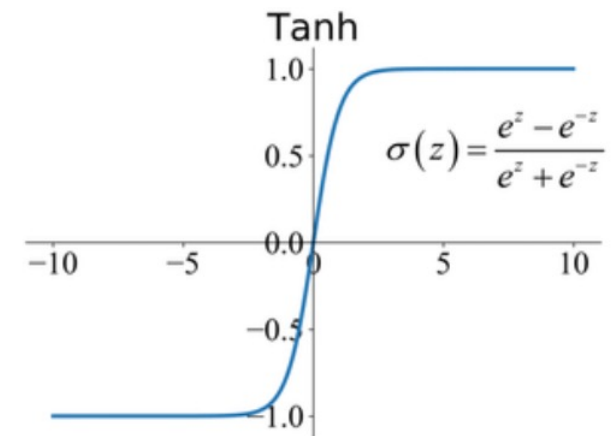
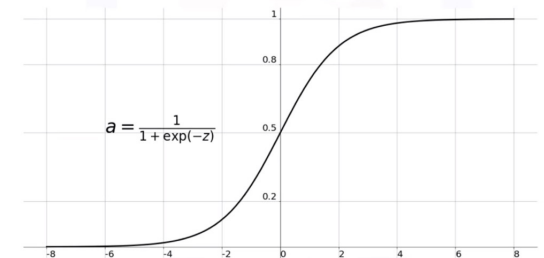
Action Potential is a fixed magnitude

Should we constrain outputs?



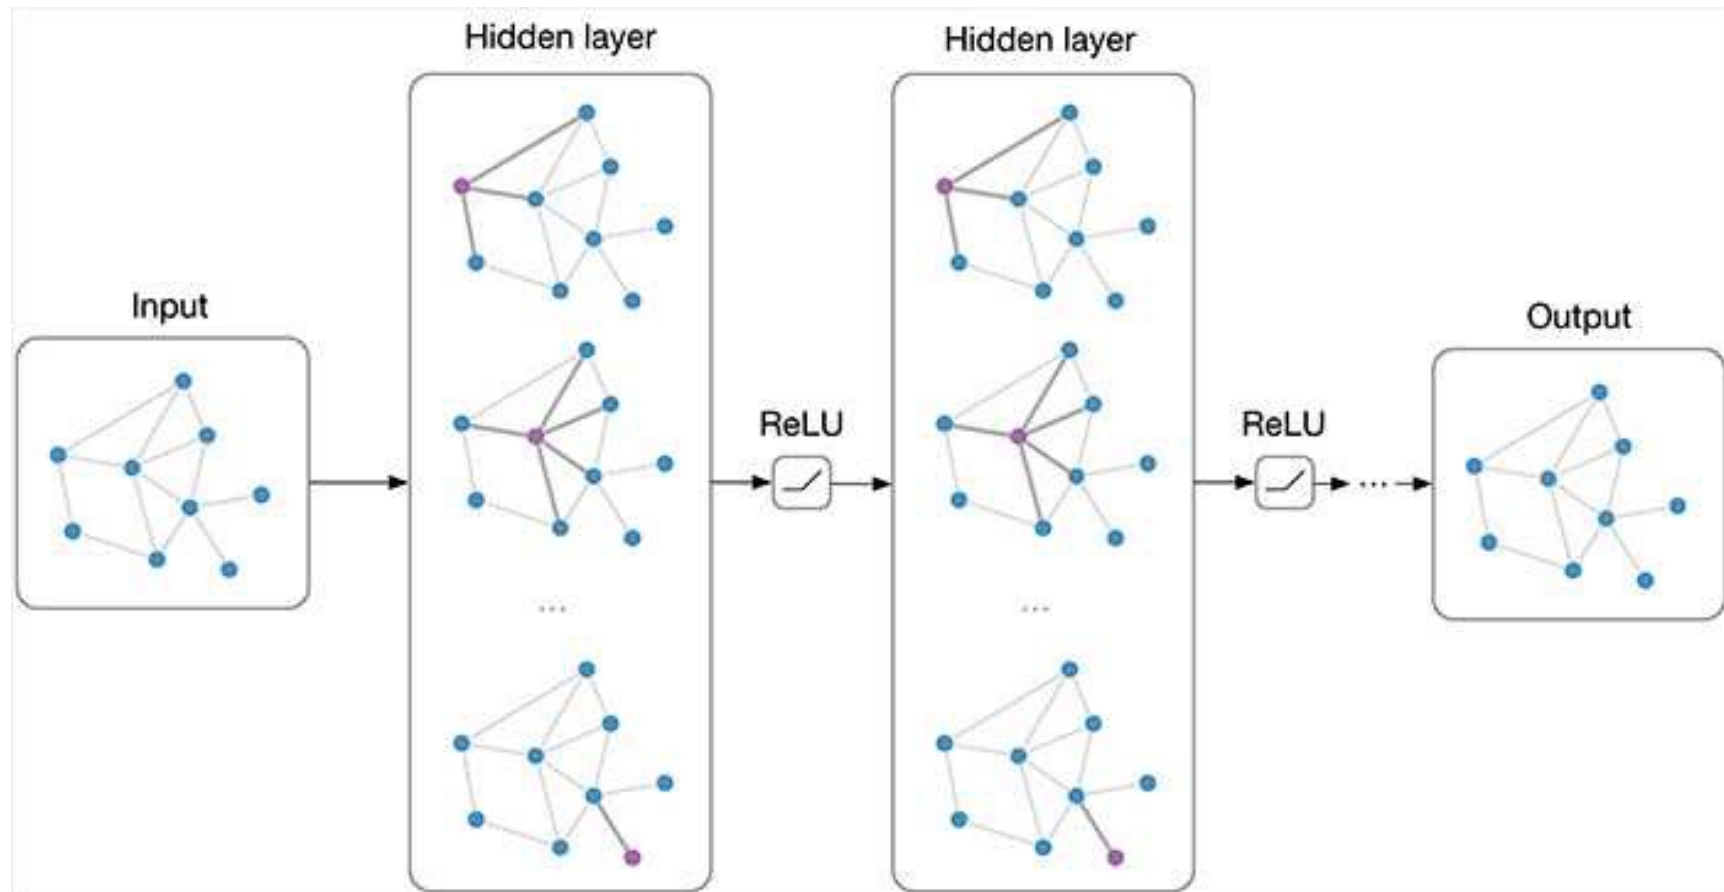
<https://en.wikipedia.org/wiki/Neuron>

Sigmoid Function



Saturated gradients:
How to solve?

Neurons can communicate in clusters: Directed Graph Neural Networks?



Questions to Ponder

- Why restrict the input of each logic gate to just 2 inputs? Why not up to n inputs?
- Should the logic gate connections be pseudo-randomly initialized, or should there be a fixed set of initial connections?
- Why not keep the logic gates fixed throughout rather than having it being selected via categorical distribution?
- Can we do with fewer operators rather than all 16 gates?
- **Can we do a similar implementation on a neural network, but restrict output to magnitude 1 for all nodes?**