

What's in a good representation?

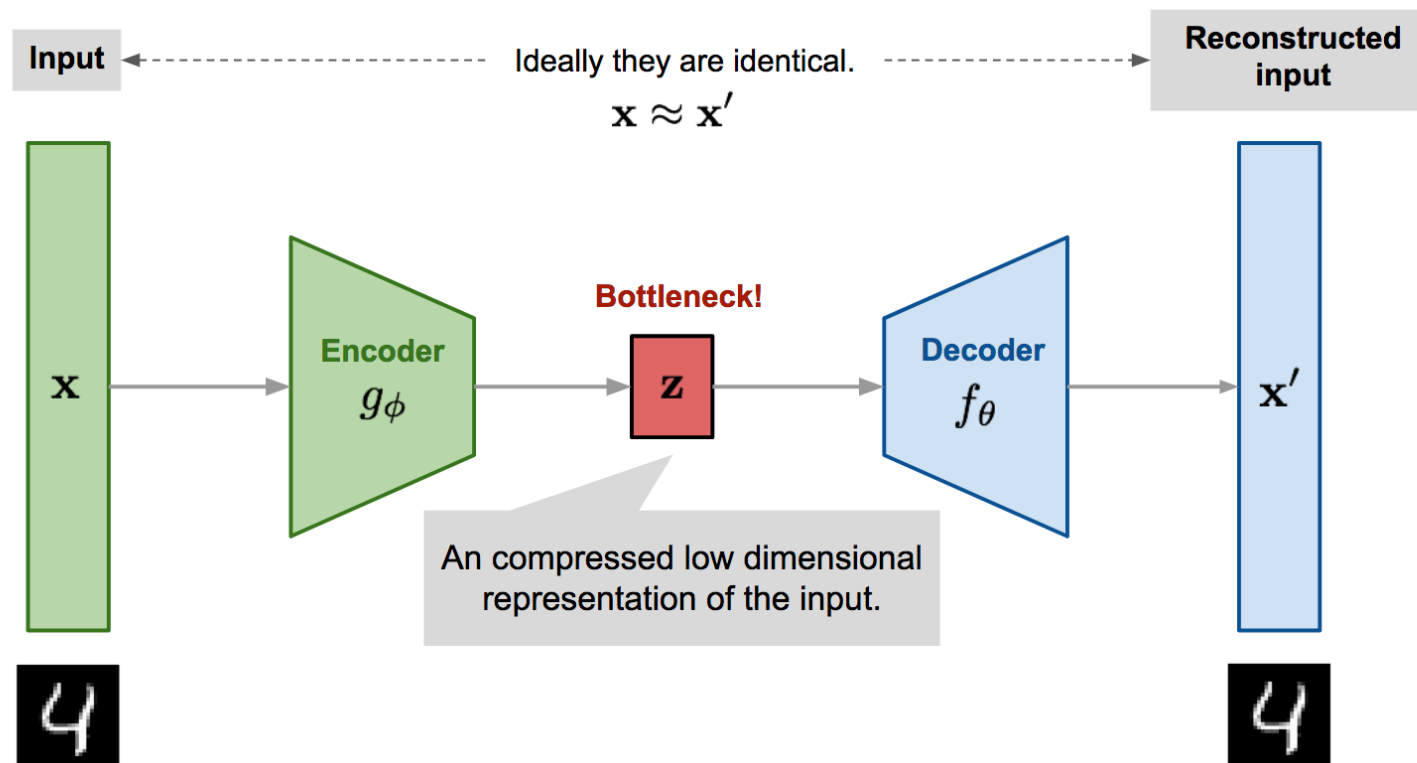
John Tan Chong Min

Aim

- Find a way to encode a suitable representation to perform decision making
- Such representation can also be how we store memories for use in the future

Autoencoders: Representation via Reconstruction

- Prioritises output clarity – may not disentangle well in latent space



Do you need to predict everything?



Transformers: Representation via Prediction

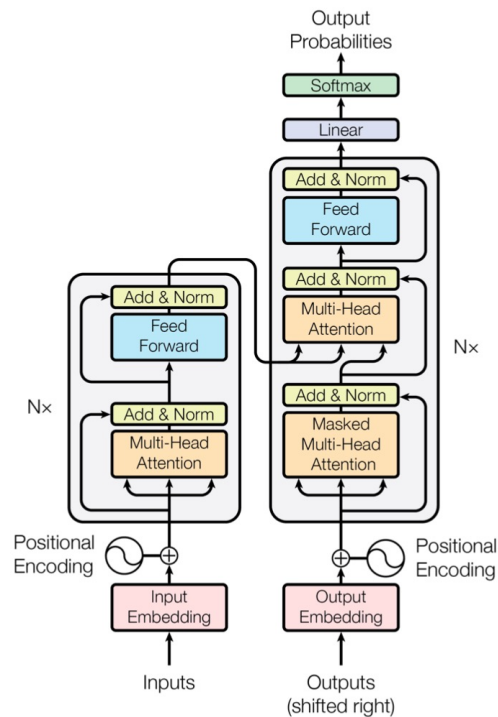
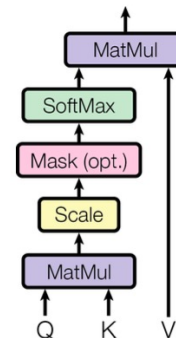


Figure 1: The Transformer - model architecture.

Scaled Dot-Product Attention



Multi-Head Attention

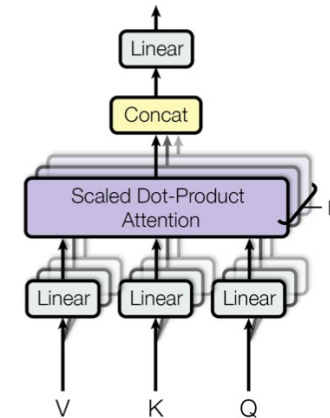


Figure 2: (left) Scaled Dot-Product Attention. (right) Multi-Head Attention consists of several attention layers running in parallel.

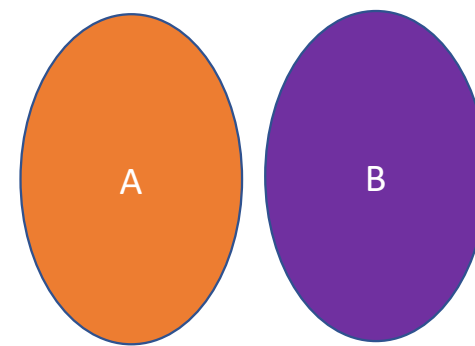
$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Taken from: Attention is all you need. Vaswani et al. (2017)

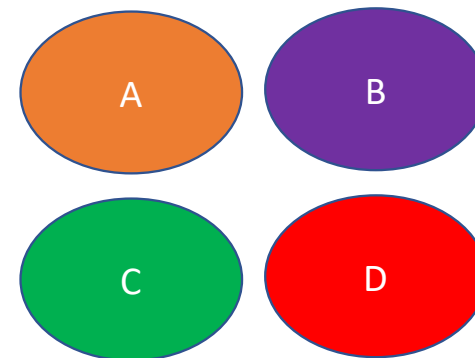
Large Self-Supervised Learning

- Self-supervised learning helps to learn better manifolds across large data
- Can work zero-shot on a new sample

Manifold of having only 2 classes

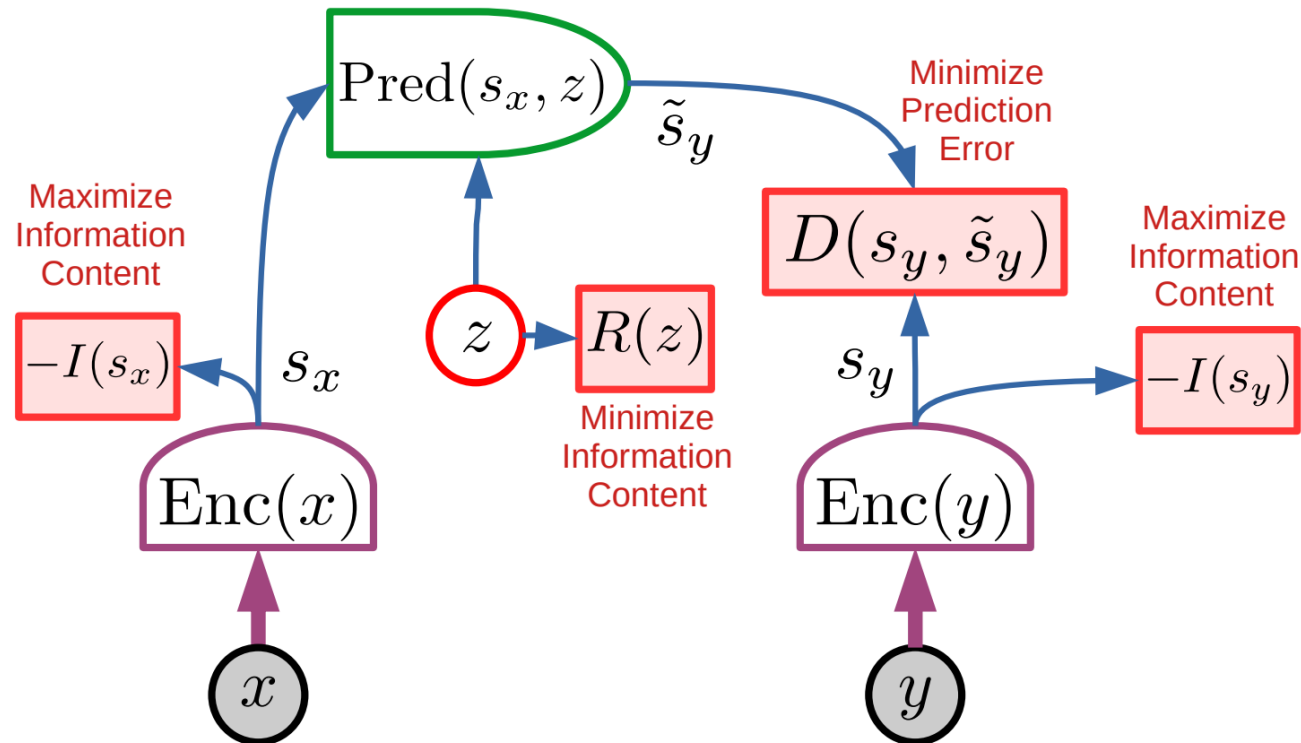


Manifold of having 4 classes



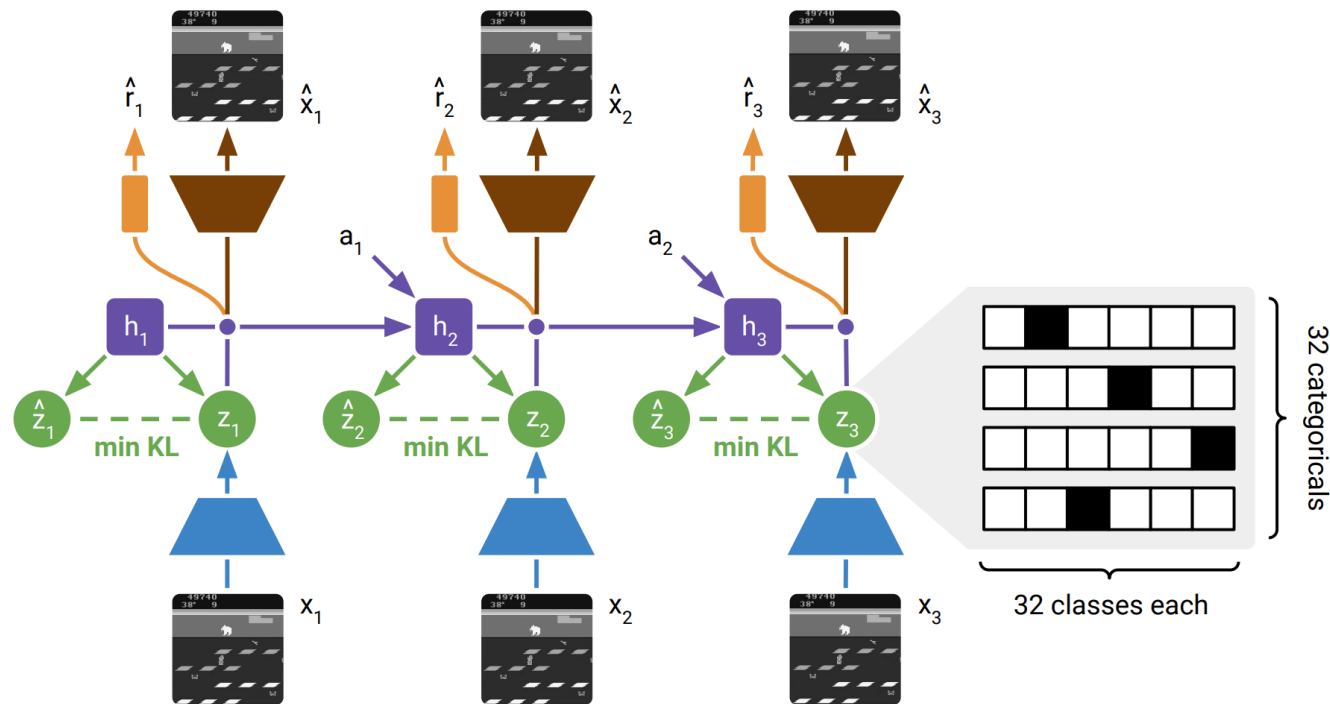
JEPA - Only use whatever is necessary to predict

- Prediction is done in latent space



World Modelling


- Use hidden representations for prediction
- Latent space is **discrete** using latent space of categorical variables



Mastering Atari with Discrete World Models. Hafner et al. 2022.


Natural Fixed Biases: Faster learning by constraints

- Sound: frequency in cilia

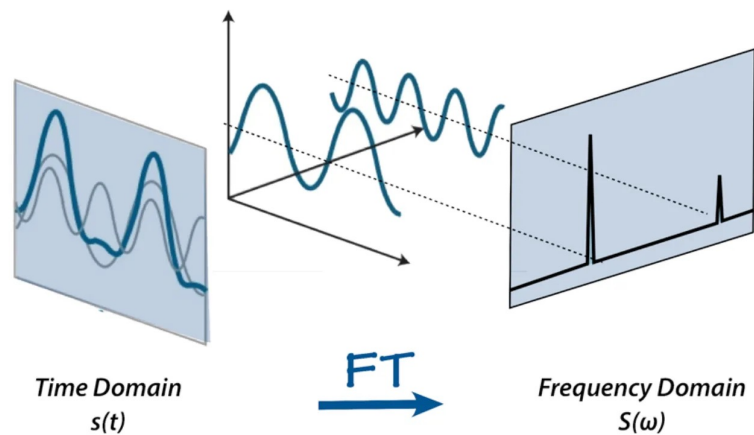


- Vision:

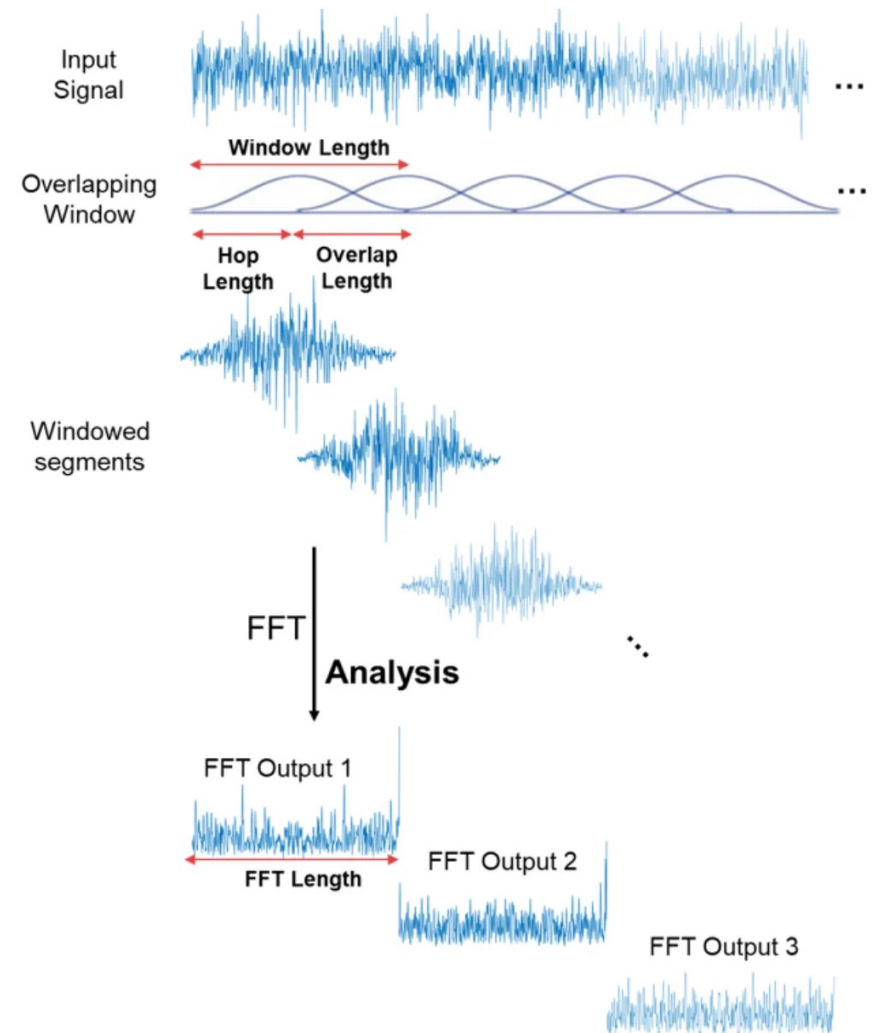
- Local patches
- Cones for Red, Green, Blue
- Rods for Black and White



Audio – Freq Modelling



Audio signal converted into frequencies



Use overlapping window to model waveform over time

Images taken from: <https://medium.com/analytics-vidhya/understanding-the-mel-spectrogram-fca2afa2ce53>

Vision – Pixel Proximity

Vertical edge detection


$$\begin{array}{|c|c|c|c|c|c|} \hline 10 & 10 & 10 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 0 & 0 & 0 \\ \hline \end{array} \quad * \quad \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline \end{array} \quad = \quad \begin{array}{|c|c|c|c|} \hline 0 & 30 & 30 & 0 \\ \hline 0 & 30 & 30 & 0 \\ \hline 0 & 30 & 30 & 0 \\ \hline 0 & 30 & 30 & 0 \\ \hline \end{array}$$

6x6

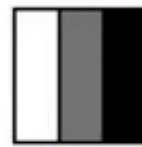
3x3

4x4


↑



*

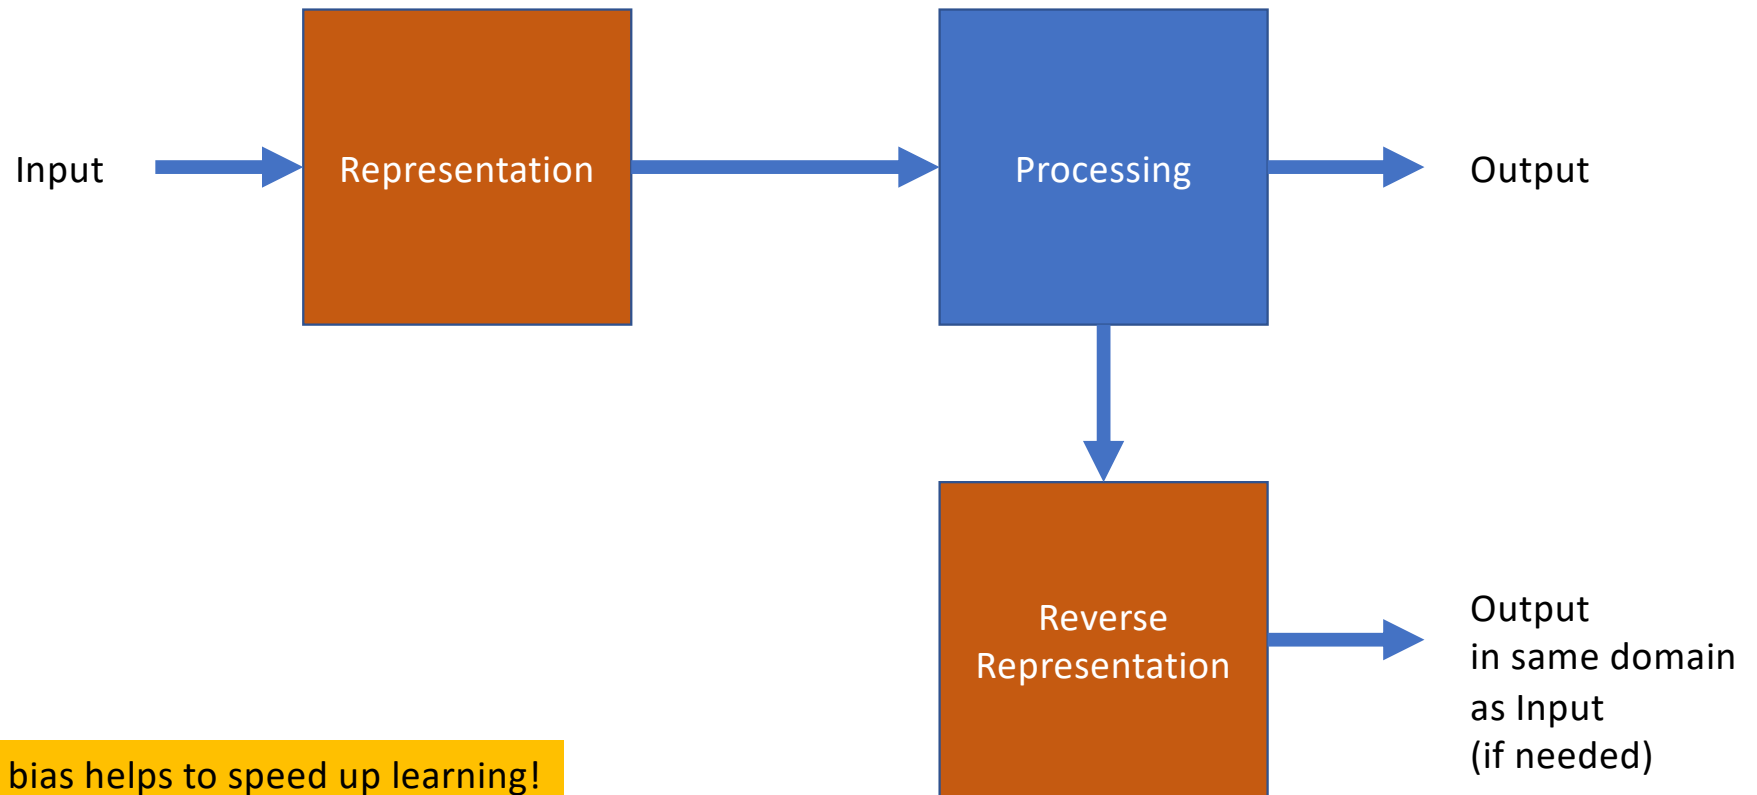


↑ ↑ ↑



Andrew Ng

Information Pipeline – Bias for Representation




The next level


Hierarchical Prediction

Hierarchical Prediction is the future

- Hierarchical prediction of more than just next token, but broader prediction at higher levels
- Higher level prediction can be more abstract and less detailed than lower levels



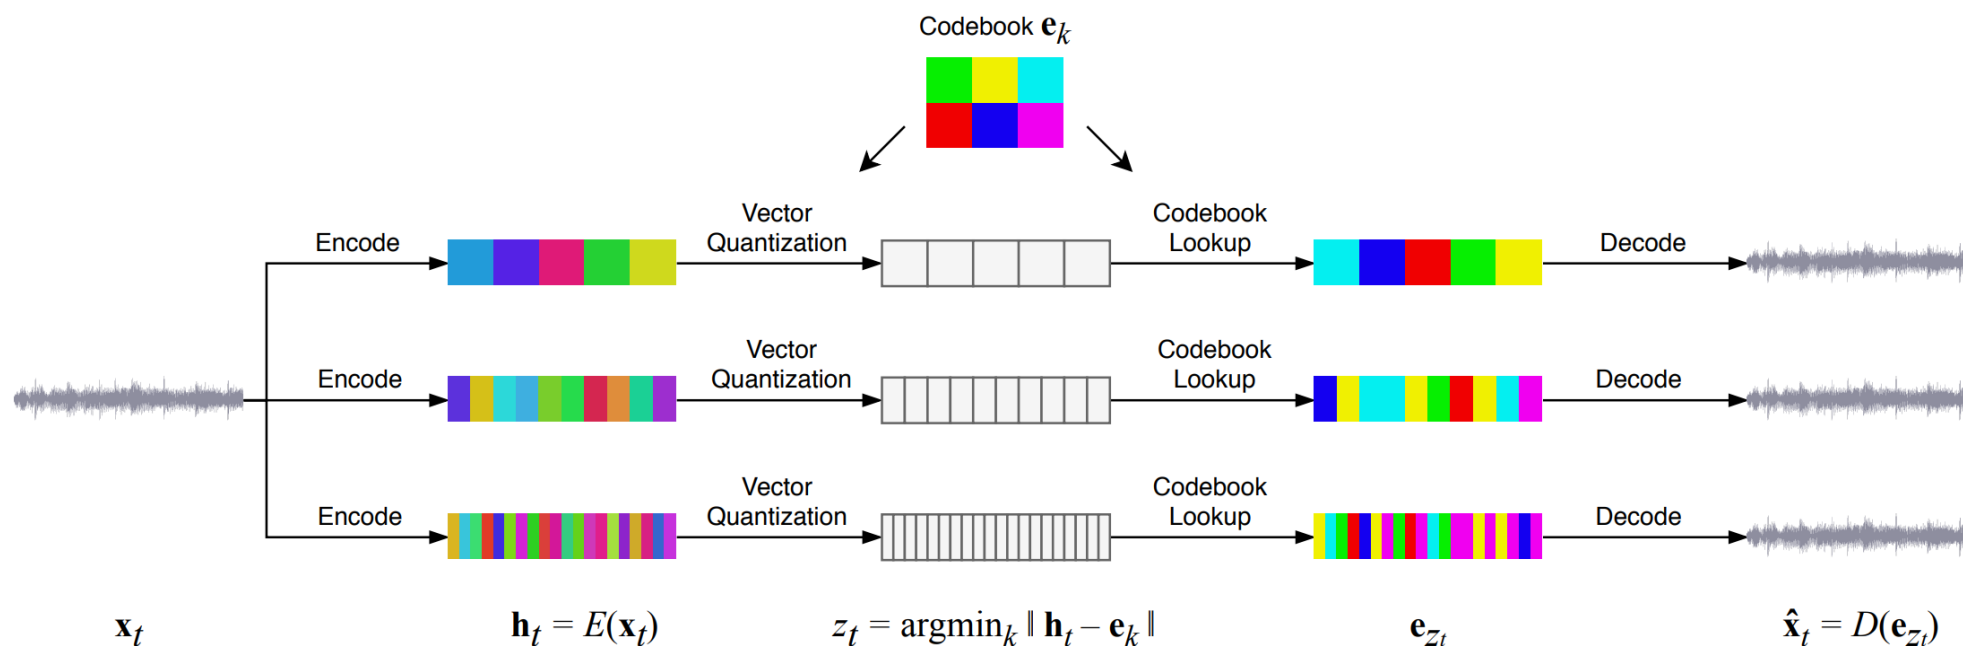
Evidence of a predictive coding hierarchy in the human brain listening to speech.
Caucheteux. 2022. Nature Human Behaviour.



How to represent part-whole hierarchies in a neural network. Hinton. 2021.

Hierarchical Prediction – Jukebox (OpenAI)


- Hierarchical Quantization of sound from coarse-grained input to fine-grained layers
- Codebook entries are **finite**



Jukebox: A Generative Model for Music. Dhariwal et. al. 2020.

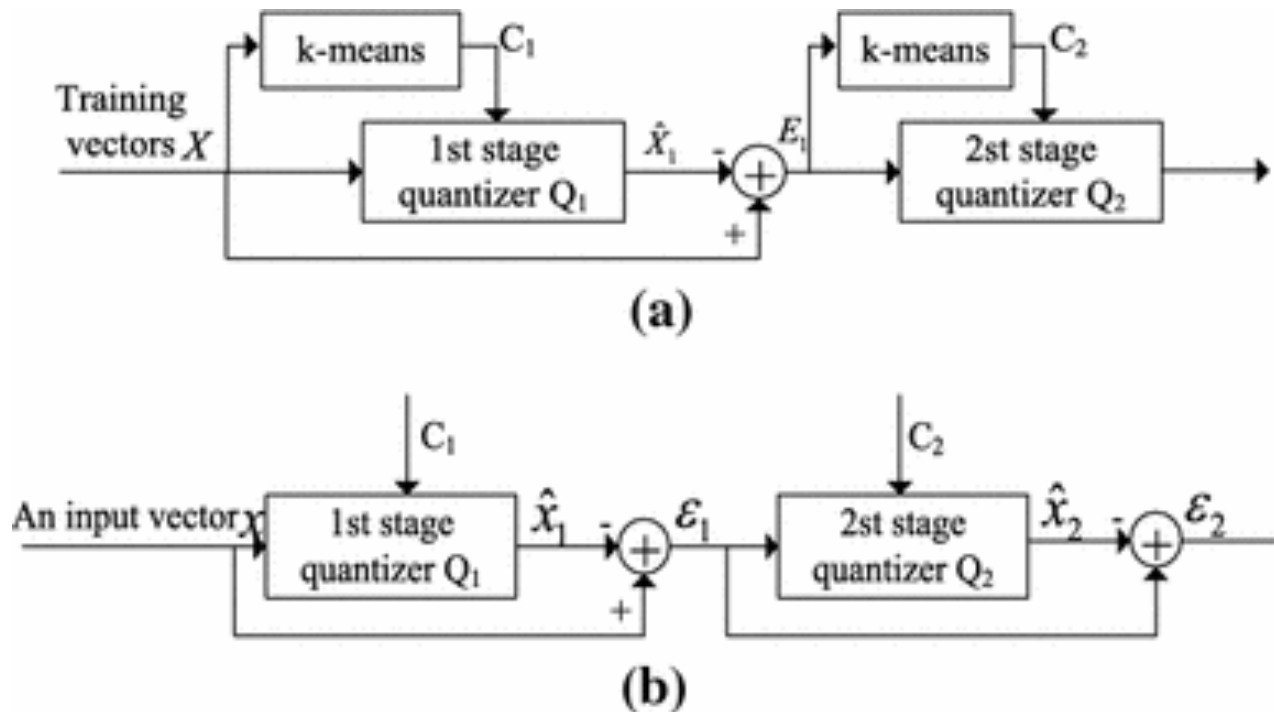
Hierarchical Prediction – Jukebox (OpenAI)

- Conditional generation of coarse-grained input to fine-grained layers



Residual Vector Quantization


- Hierarchical Quantization from coarse-grained input to fine-grained layers
- Akin to compositionality?




Optimized residual vector quantization for efficient approximate nearest neighbor search. Ai et al. 2015.

Hierarchical Prediction - Feature Pyramid Network


- Hierarchical prediction from coarse-grained image to fine-grained image




(a) Featurized image pyramid



(b) Single feature map



(c) Pyramidal feature hierarchy

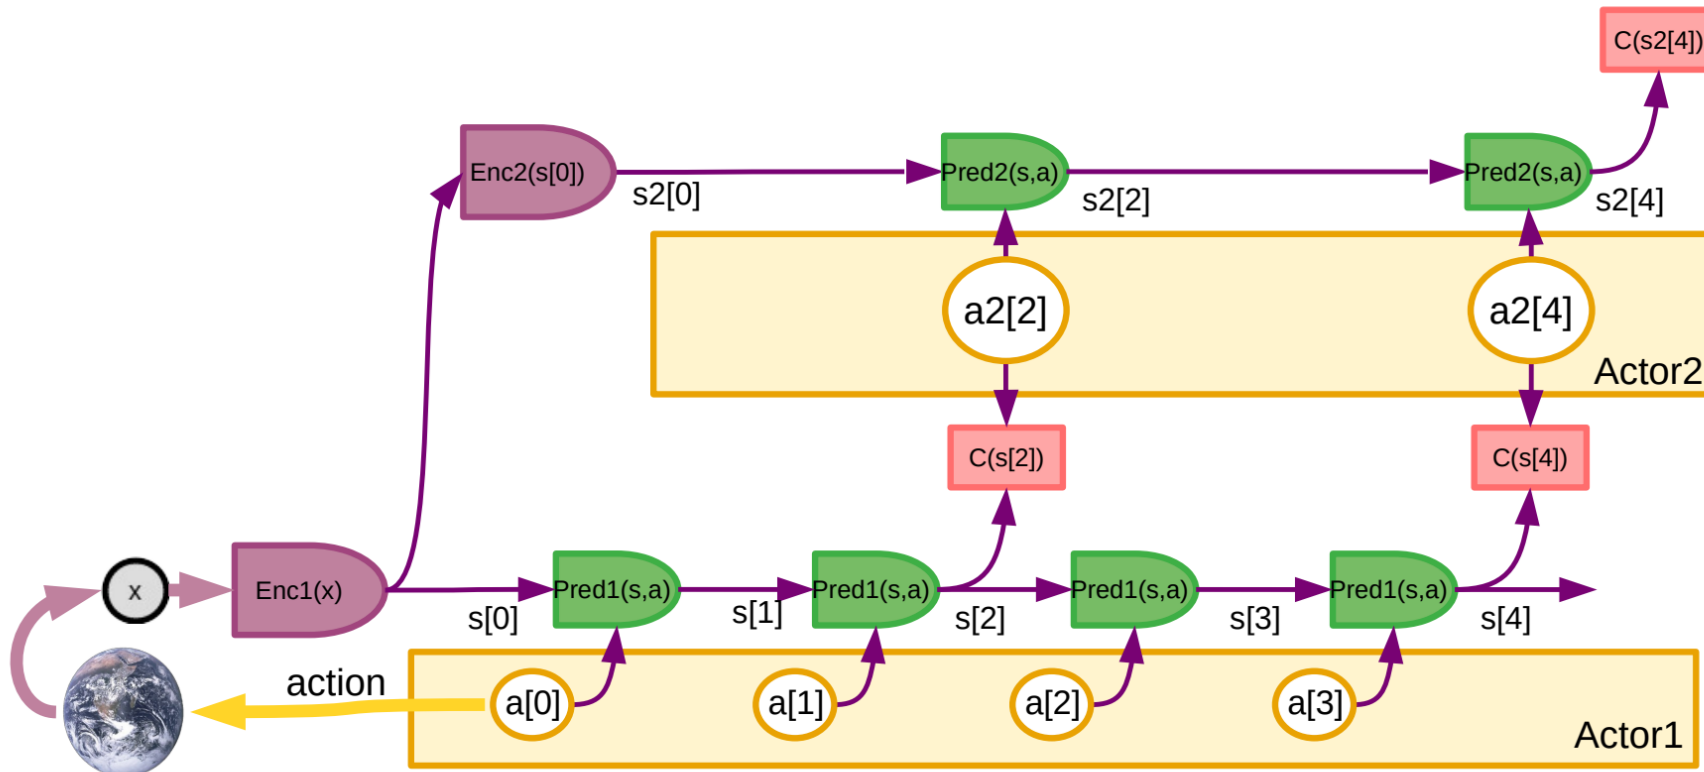


(d) Feature Pyramid Network

Feature Pyramid Networks for Object Detection. Lin et al. 2017.

Hierarchical JEPA

- Hierarchical prediction of actions from the highest level action to the lowest level action



A Path towards Autonomous Machine Intelligence. Yann LeCun. 2022.

Hierarchical Action Prediction

- Hierarchical prompting of actions from broad action to specific actions

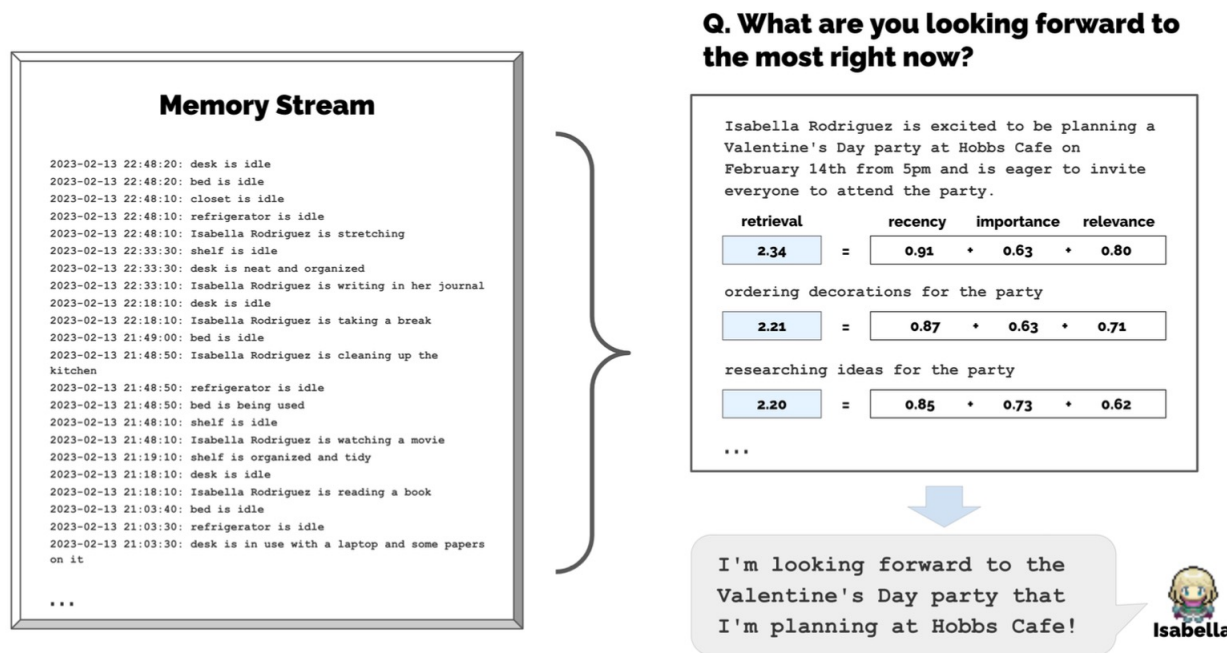


Figure 6: The memory stream comprises a large number of observations that are relevant and irrelevant to the agent's current situation. Retrieval identifies a subset of these observations that should be passed to the language model to condition its response to the situation.

Transformers:

Can a Transformer perform hierarchical generation?

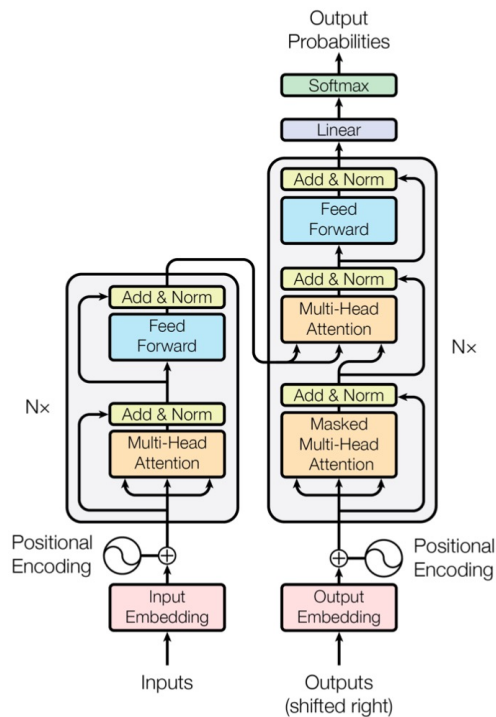
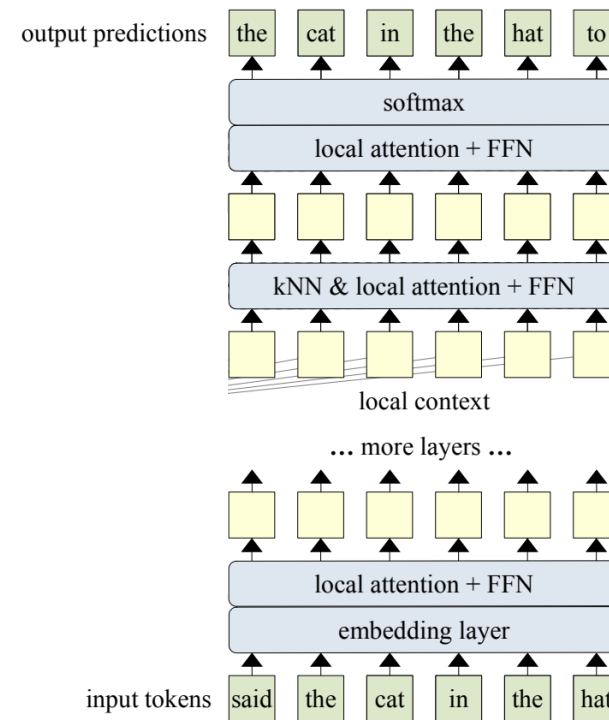


Figure 1: The Transformer - model architecture.

Attention is all you need. Vaswani et al. (2017)



Memorizing Transformers. Wu et al. 2022.

Questions to Ponder

- Do we learn from experience, or from natural fixed bias? Or both?
- Should we do prediction whereby we map back to input space (like tokens in Transformers), or should we just predict the latent space? What are the benefits and drawbacks?
- Should we use hierarchical generation? Is our brain hierarchical or more flat like what Jeff Hawkins proposes in “Thousand Brains Theory”?
- Should we represent latent space as continuous or discrete? Would an unbounded length of discrete tokens be sufficient to represent continuous spaces?